

Introduction

Team: Solution Accepted;

Members: Junxian Chen, Wen Chia Yang, Zihua Weng;

[Realm](#) is a mobile database that runs directly inside phones, tablets or wearables.

Feature 1: Adding more query methods

Background

Queries are the most common functions of a database.

Realm's query engine uses a Fluent interface to construct multi-clause queries.

```
public class User extends RealmObject {

    @PrimaryKey
    private String      name;
    private int         age;

    @Ignore
    private int         sessionId;

    // Standard getters & setters generated by your IDE...
    public String getName() { return name; }
    public void setName(String name) { this.name = name; }
    public int getAge() { return age; }
    public void setAge(int age) { this.age = age; }
    public int getSessionId() { return sessionId; }
    public void setSessionId(int sessionId) { this.sessionId = sessionId; }
}
```

Let say if we want to find all users named John or Peter, we would write:

```
// Build the query looking at all users:
RealmQuery<User> query = realm.where(User.class);

// Add query conditions:
query.equalTo("name", "John");
```

```

query.or().equalTo("name", "Peter");

// Execute the query:
RealmResults<User> result1 = query.findAll();

// Or alternatively do the same all at once (the "Fluent interface"):
RealmResults<User> result2 = realm.where(User.class)
    .equalTo("name", "John")
    .or()
    .equalTo("name", "Peter")
    .findAll();

```

Here Realm has already implemented a bundle of query functions for us, such as `where()`, `equalTo()`, `or()`, `greaterThan()`, `count()`, etc. Say, we would like to add these useful functions into the Realm Query module, then the following research will be helpful to the developers who intend to realize the features.

Research

As we can see from the above example, Realm first creates a `RealmQuery` object and then adds query conditions on it.

Take `equalTo()` as an example, when `equalTo()` in `io/realm/RealmQuery.java` is executed, `RealmQuery` calls `equalTo()` function from `io/realm/internal/TableQuery.java`.

```

public RealmQuery<E> equalTo(String fieldName, @Nullable String value) {
    return this.equalTo(fieldName, value, Case.SENSITIVE);
}

public RealmQuery<E> equalTo(String fieldName, @Nullable String value, Case casing) {
    realm.checkIfValid();

    return equalToWithoutThreadValidation(fieldName, value, casing);
}

private RealmQuery<E> equalToWithoutThreadValidation(String fieldName, @Nullable String
value, Case casing) {
    FieldDescriptor fd = schema.getColumnIndices(fieldName, RealmFieldType.STRING);
    this.query.equalTo(fd.getColumnIndices(), fd.getNativeTablePointers(), value, casing);
    return this;
}

```

Then `equalTo()` will call `nativeEqual` which is a Realm C++ library method and return a `TableQuery` object itself.

`TableQuery` class supports all low level methods (`equalTo/greaterThan/lessThan/count`) a table has. All mentioned native communications to the Realm C++ library are also handled by this class.

```
// Equals
public TableQuery equalTo(long[] columnIndexes, long[] tablePtrs, @Nullable String value,
    Case caseSensitive) {
    nativeEqual(nativePtr, columnIndexes, tablePtrs, value, caseSensitive.getValue());
    queryValidated = false;
    return this;
}
```

```
private native void nativeEqual(long nativeQueryPtr, long[] columnIndex,
    long[] tablePtrs, long value);
```

Going back to the `equalTo()` function in `RealmQuery`, the `query` field in `RealmQuery` is assigned to the return `TableQuery` object. Finally, the function returns a `RealmQuery` object.

Conclusion

In conclusion, for developers who want to add more query function to Realm, one should add customized method to `io/realm/RealmQuery.java` and the corresponding method in `io/realm/internal/TableQuery.java` which calls a defined C++ library function.

RealmQuery		
createQuery(Realm, Class<E>)		RealmQuery<E>
createDynamicQuery(DynamicRealm, String)		RealmQuery<E>
createQueryFromResult(RealmResults<E>)		RealmQuery<E>
createQueryFromList(RealmList<E>)		RealmQuery<E>
isClassForRealmModel(Class<?>)		boolean
isValid()		boolean
isNull(String)		RealmQuery<E>
isNotNull(String)		RealmQuery<E>
equalTo(String, String)		RealmQuery<E>
equalTo(String, String, Case)		RealmQuery<E>
equalToWithoutThreadValidation(String, String, Case)		RealmQuery<E>
equalTo(String, Byte)		RealmQuery<E>
equalToWithoutThreadValidation(String, Byte)		RealmQuery<E>
equalTo(String, byte[])		RealmQuery<E>
equalTo(String, Short)		RealmQuery<E>
equalToWithoutThreadValidation(String, Short)		RealmQuery<E>

NativeObject		
getNativePtr()	long	
getNativeFinalizerPtr()	long	

TableQuery		
getNativePtr()	long	
getNativeFinalizerPtr()	long	
getTable()	Table	
validateQuery()	void	
group()	TableQuery	
endGroup()	TableQuery	
or()	TableQuery	
not()	TableQuery	
equalTo(long[], long[], long)	TableQuery	
notEqualTo(long[], long[], long)	TableQuery	
greaterThan(long[], long[], long)	TableQuery	
greaterThanOrEqualTo(long[], long[], long)	TableQuery	
lessThan(long[], long[], long)	TableQuery	

Feature 2: Adding more encryption methods

Background

From [the official documentation](#) we know that the Realm database file can be encrypted with a 512-bit key (i.e. 64 bytes), using standard AES-256 encryption algorithm. In the future we may want to add more encryption methods like [Triple DES](#), [RSA](#), [ChaCha20](#) or [TwoFish](#) to enhance security of the database. To do so, we have to find out where these encryption methods are adopted.

Research

First let's start with a code snippet which is likely to appear in our Android project:

```
byte[] key = new byte[64];
new SecureRandom().nextBytes(key);
RealmConfiguration config = new RealmConfiguration.Builder()
    .encryptionKey(key)
    .build();

Realm realm = Realm.getInstance(config);
```

The snippet above randomly generated a 64-byte key which served as the encryption key of our Realm database. Here a new `RealmConfiguration.Builder` was created and the method `encryptionKey(key)` was called. `RealmConfiguration.Builder` is an inner class of `RealmConfiguration` so we should check inside `RealmConfiguration.java`:

RealmConfiguration.java:548

```
public Builder encryptionKey(byte[] key) {
    //noinspection ConstantConditions
    if (key == null) {
        throw new IllegalArgumentException("A non-null key must be provided");
    }
    if (key.length != KEY_LENGTH) {
        throw new IllegalArgumentException(String.format(Locale.US,
            "The provided key must be %s bytes. Yours was: %s",
            KEY_LENGTH, key.length));
    }
    this.key = Arrays.copyOf(key, key.length);
}
```

```
    return this;
}
```

We navigated to `RealmConfiguration.java`, and inside the `RealmConfiguration.Builder` class we located the `encryptionKey(key)` method. Basically what it did was assigning the key value to the Builder and then return the Builder itself for chaining.

RealmConfiguration.java:794

```
public RealmConfiguration build() {
    // Check that readOnly() was applied to legal configuration. Right now it
    // should only be allowed if
    // an assetFile is configured
    if (readOnly) {
        // ...

    return new RealmConfiguration(directory,
        fileName,
        getCanonicalPath(new File(directory, fileName)),
        assetFilePath,
        key,
        schemaVersion,
        migration,
        deleteRealmIfMigrationNeeded,
        durability,
        createSchemaMediator(modules, debugSchema),
        rxFactory,
        initialDataTransaction,
        readOnly,
        compactOnLaunch,
        false
    );
}
```

Finally by calling the `build()` method, a `RealmConfiguration` was created. Please note that the key is now passed to the new `RealmConfiguration`. Now we are wondering how the key was used across the Realm project. Inside `RealmConfiguration`, the only way to access the key is calling `getEncryptionKey()`:

RealmConfiguration.java:152

```
public byte[] getEncryptionKey() {
```

```
    return key == null ? null : Arrays.copyOf(key, key.length);
}
```

However, there are no usages of `getEncryptionKey()` across the project. We should look somewhere else.

Then we noticed “`Realm realm = Realm.getInstance(config)`”. The configuration was passed to a static method inside the Realm class called “`getInstance()`”. We traced the passing route of the configuration:

Realm.java:410

```
public static Realm getInstance(RealmConfiguration configuration) {
    //noinspection ConstantConditions
    if (configuration == null) {
        throw new IllegalArgumentException(NULL_CONFIG_MSG);
    }
    return RealmCache.createRealmOrGetFromCache(configuration, Realm.class);
}
```

RealmCache:286

```
static <E extends BaseRealm> E createRealmOrGetFromCache(RealmConfiguration
configuration,
    Class<E> realmClass) {
    RealmCache cache = getCache(configuration.getPath(), true);

    return cache.doCreateRealmOrGetFromCache(configuration, realmClass);
}
```

RealmCache:293

```
private synchronized <E extends BaseRealm> E
doCreateRealmOrGetFromCache(RealmConfiguration configuration,
    Class<E> realmClass)
```

RealmCache:312 (Inside RealmCache.**doCreateRealmOrGetFromCache()**)

```
OsRealmConfig osConfig = new OsRealmConfig.Builder(configuration).build();
```

OsRealmConfig.java:87 (The declaration of OsRealmConfig.Builder)

```
public static class Builder
```

OsRealmConfig.java:98 (The constructor of OsRealmConfig.Builder)

```
public Builder(RealmConfiguration configuration) {  
    this.configuration = configuration;  
}
```

OsRealmConfig.java:150 (The build() method of OsRealmConfig.Builder)

```
public OsRealmConfig build() {  
    return new OsRealmConfig(configuration, fifoFallbackDir,  
        autoUpdateNotification, schemaInfo,  
        migrationCallback, initializationCallback);  
}
```

OsRealmConfig.java:198 (The constructor of **OsRealmConfig**)

```
private OsRealmConfig(final RealmConfiguration config,  
    String fifoFallbackDir,  
    boolean autoUpdateNotification,  
    @Nullable OsSchemaInfo schemaInfo,  
    @Nullable OsSharedRealm.MigrationCallback  
migrationCallback,  
    @Nullable OsSharedRealm.InitializationCallback  
initializationCallback)
```

OsRealmConfig.java:237 (Inside the constructor of **OsRealmConfig**)

```
// Set encryption key  
byte[] key = config.getEncryptionKey();  
if (key != null) {  
    nativeSetEncryptionKey(nativePtr, key);  
}
```

OsRealmConfig.java:368 (The declaration of **nativeSetEncryptionKey**)

```
private static native void nativeSetEncryptionKey(long nativePtr, byte[] key);
```


io_realm_internal_OsRealmConfig.cpp:97

```

JNIEXPORT void JNICALL
Java_io_realm_internal_OsRealmConfig_nativeSetEncryptionKey(JNIEnv* env,
jclass,
jlong native_ptr,
jbyteArray j_key_array)
{
    TR_ENTER_PTR(native_ptr)
    try {
        JByteArrayAccessor jarray_accessor(env, j_key_array);
        auto& config = *reinterpret_cast<Realm::Config*>(native_ptr);
        // Encryption key should be set before creating sync_config.
        REALM_ASSERT(!config.sync_config);
        config.encryption_key = jarray_accessor.transform<std::vector<char>>>();
    }
    CATCH_STD()
}

```

Here, we have found that the encryption key setter is a native C++ function. We continued our search but did not find anything related to encryption in our [realm-java project](#). Later we searched the Internet and found [a page on StackOverflow](#), which said the encryption module was embedded in the [realm-core project](#). The default encryption algorithm in Realm is AES-256 which is declared in [aes_cryptor.hpp](#). Realm handles different encryption algorithms by `realm::util::encryption_read_barrier` and `realm::util::encryption_write_barrier` using different file mapping ([alloc_slab.cpp](#)).

```

realm::util::encryption_read_barrier(addr, Array::header_size,
map->get_encrypted_mapping(),Array::get_byte_size_from_header);

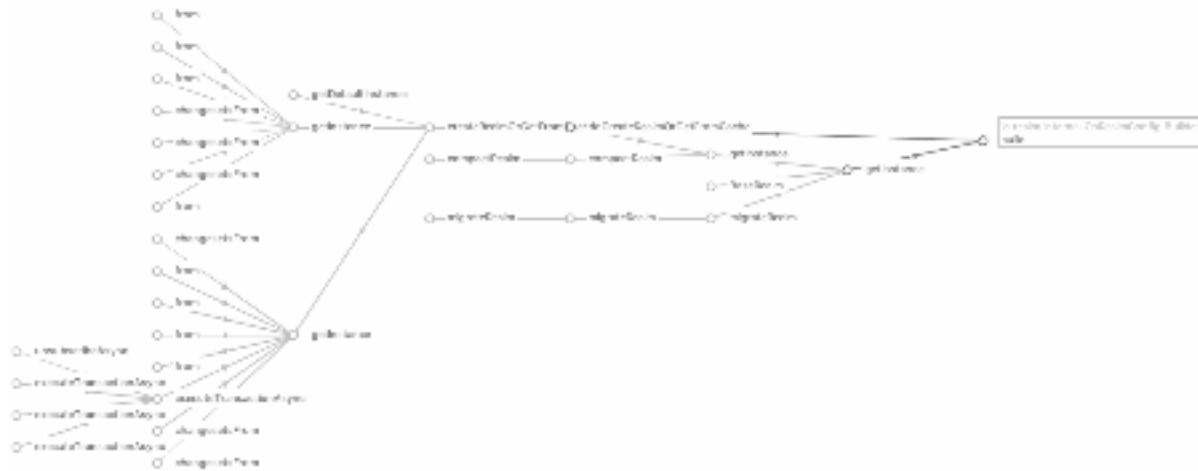
```

Here the encrypted_mapping is retrieved from [encrypted_file_mapping.cpp](#).

Theoretically, in [encrypted_file_mapping.cpp](#), we could add a new Cryptor class here and implement different encryption methods.

Conclusion

OsRealmConfig.build() is called by multiple places. Here in our feature 2, we only focus on the path: build() -> doCreateRealmOrGetFromCache -> createRealmOrGetFromCache -> getInstance.



nativeSetSchemaConfig is the only one method called in OsRealmConfig.



By tracing the config we located the feature can be added by modifying the realm-core project.